

# Robotics 2

## AdaBoost for People and Place Detection

Kai Arras, Cyrill Stachniss,  
Maren Bennewitz, Wolfram Burgard



# Classification

- Classification algorithms are **supervised algorithms** to predict **categorical labels**
- Differs from **regression** which is a supervised technique to predict **real-valued labels**

## Formal problem statement:

- **Produce a function that maps**

$$C : \mathcal{X} \rightarrow \mathcal{Y}$$

- **Given a training set**


$$\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$$

$y \in \mathcal{Y}$       label  
 $\mathbf{x} \in \mathcal{X}$       training sample

# Classification

## Error types

		True value	
		$T$	$N$
Predicted value of the classifier	$T'$	True Positive	False Positive
	$N'$	False Negative	True Negative




**Detected**  
**Not Detected**

- **Precision** =  $TP / (TP + FP)$
- **Recall** =  $TP / (TP + FN)$

Many more measures...

# Classification


## Linear vs. Non-Linear Classifier, Margin



# Classification

## Overfitting

- Overfitting occurs when a model begins to memorize the **training data** rather than learning the underlying relationship
- Occurs typically when fitting a statistical model with too many parameters
- Overfitted models explain training data perfectly but they **do not generalize!**
- There are techniques to avoid overfitting such as regularization or cross-validation



# Boosting

- An **ensemble technique** (a.k.a. committee method)
- Supervised learning: given  $\langle \text{samples } x, \text{ labels } y \rangle$
- Learns an accurate **strong classifier** by combining an ensemble of inaccurate “rules of thumb”
- **Inaccurate rule**  $h(x_i)$ : “weak” classifier, weak learner, basis classifier, feature
- **Accurate rule**  $H(x_i)$ : “strong” classifier, final classifier
- Other **ensemble techniques** exist: Bagging, Voting, Mixture of Experts, etc.

# AdaBoost

- Most popular algorithm: **AdaBoost**

*[Freund et al. 95], [Schapire et al. 99]*

- Given an ensemble of weak classifiers  $h(x_i)$ , the combined strong classifier  $H(x_i)$  is obtained by a **weighted majority voting scheme**

$$f(x_i) = \sum_{t=1}^T \alpha_t h_t(x_i) \quad H(x_i) = \text{sgn}(f(x_i))$$

- AdaBoost in Robotics:

*[Viola et al. 01], [Treptow et al. 04], [Martínez-Mozos et al. 05], [Rottmann et al. 05], [Monteiro et al. 06], [Arras et al. 07]*

# AdaBoost

Why is AdaBoost interesting?

1. It tells you what the **best "features"** are
  2. What the **best thresholds** are, and
  3. How to **combine them to a classifier**
- 
- AdaBoost can be seen as a **principled feature selection strategy**
  - Classifier design becomes **science**, not art



# AdaBoost

- AdaBoost is a **non-linear classifier**
- Has **good generalization properties**: can be proven to maximize the margin
- Quite robust to **overfitting**
- Very **simple** to implement
- **Prerequisite:**  
weak classifier must be better than chance:  
 $\text{error} < 0.5$  in a binary classification problem

# AdaBoost

- **Possible Weak Classifiers:**

- **Decision stump:**

- Single axis-parallel partition of space

- **Decision tree:**

- Hierarchical partition of space

- **Multi-layer perceptron:**

- General non-linear function approximators

- **Support Vector Machines (SVM):**

- Linear classifier with RBF Kernel


- Trade-off between diversity among weak learners versus their accuracy. Can be complex, see literature
- Decision stumps are **a popular choice**

# AdaBoost: Weak Classifier

## Decision stump

- Simple-most type of **decision tree**
- Equivalent to linear classifier defined by affine hyperplane
- Hyperplane is orthogonal to axis with which it intersects in threshold  $\theta$
- Commonly not used on its own
- Formally,

$$h(x; j, \theta) = \begin{cases} +1 & x_j > \theta \\ -1 & \text{else} \end{cases}$$



where  $x$  is ( $d$ -dim.) training sample,  $j$  is dimension


# AdaBoost: Weak Classifier

- **Train a decision stump on weighted data**

$$(j^*, \theta^*) = \operatorname{argmin}_{j, \theta} \left\{ \sum_{i=1}^n w_t(i) \mathbf{I}(y_i \neq h_t(x_i)) \right\}$$

- **This consists in...**

Finding an optimum parameter  $\theta^*$  for each dimension  $j = 1 \dots d$  and then select the  $j^*$  for which the weighted error is minimal.



# AdaBoost: Weak Classifier

## A simple training algorithm for stumps:

$\forall j = 1 \dots d$

Sort samples  $x_i$  in ascending order along dimension  $j$

$\forall i = 1 \dots n$

Compute  $n$  cumulative sums  $w_{cum}^j(i) = \sum_{k=1}^i w_k y_k$

end

Threshold  $\theta_j$  is at extremum of  $w_{cum}^j$

Sign of extremum gives direction  $p_j$  of inequality



end

Global extremum in all  $d$  sums  $w_{cum}$  gives  
**threshold**  $\theta^*$  and **dimension**  $j^*$


# AdaBoost: Weak Classifier

## Training algorithm for stumps: Intuition

- **Label  $y$  :**  
red: +  
blue: -
- Assuming all weights = 1



$$w_{cum}^j(i) = \sum_{k=1}^i w_k y_k$$




# AdaBoost: Algorithm

Given the **training data**  $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$   $\mathbf{x} \in \mathcal{X}$   $y \in \mathcal{Y}$

1. Initialize weights  $w_t(i) = 1/n$
2. For  $t = 1, \dots, T$ 
  - Train a **weak classifier**  $h_t(x)$  on weighted training data minimizing the error
$$\varepsilon_t = \sum_{i=1}^n w_t(i) \mathbf{I}(y_i \neq h_t(x_i))$$
  - Compute voting weight of  $h_t(x)$ : 
$$\alpha_t = \frac{1}{2} \log((1 - \varepsilon_t)/\varepsilon_t)$$
  - Recompute weights: 
$$w_{t+1}(i) = w_t(i) \exp\{-\alpha_t y_i h_t(x_i)\} / Z_t$$
3. Make predictions using the final **strong classifier**

# AdaBoost: Voting Weight

- Computing the **voting weight**  $\alpha_t$  of a weak classifier
- $\alpha_t$  measures the **importance** assigned to  $h_t(x_i)$





# AdaBoost: Weight Update

- Looking at the weight update step:

$$w_{t+1}(i) = w_t(i) \exp\{-\alpha_t y_i h_t(x_i)\} / Z_t$$

**Normalizer** such  
 $Z_t$  : that  $w_{t+1}$  is a prob.  
distribution

$$\exp\{-\alpha_t y_i h_t(x_i)\} = \begin{cases} < 1, & y_i = h_t(x_i) \\ > 1, & y_i \neq h_t(x_i) \end{cases}$$

- Weights of misclassified training samples are **increased**
- Weights of correctly classified samples are **decreased**
- Algorithm generates weak classifier by training the next learner on the **mistakes of the previous one**
- Now we understand the name: **AdaBoost** comes from **adaptive Boosting**

# AdaBoost: Strong Classifier

- **Training is completed...**

The weak classifiers  $h_{1...T}(x)$  and their voting weight  $\alpha_{1...T}$  are now fix

- **The resulting strong classifier is**

The diagram shows the equation for the strong classifier  $H(x_i)$  inside a rounded rectangle. An arrow points from the text 'Put your data here' to the variable  $x_i$  in the equation. To the right of the rectangle, an arrow points to the text 'Class Result  $\{+1, -1\}$ '.

$$H(x_i) = \text{sgn} \left( \sum_{t=1}^T \alpha_t h_t(x_i) \right) \longrightarrow \text{Class Result } \{+1, -1\}$$

Put your data here

Weighted majority voting scheme


# AdaBoost: Algorithm

Given the **training data**  $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$   $\mathbf{x} \in \mathcal{X}$   $y \in \mathcal{Y}$

1. Initialize weights  $w_t(i) = 1/n$
2. For  $t = 1, \dots, T$ 
  - Train a **weak classifier**  $h_t(x)$  on weighted training data minimizing the error
$$\varepsilon_t = \sum_{i=1}^n w_t(i) \mathbf{I}(y_i \neq h_t(x_i))$$
  - Compute voting weight of  $h_t(x)$ : 
$$\alpha_t = \frac{1}{2} \log((1 - \varepsilon_t)/\varepsilon_t)$$
  - Recompute weights: 
$$w_{t+1}(i) = w_t(i) \exp\{-\alpha_t y_i h_t(x_i)\} / Z_t$$
3. Make predictions using the final **strong classifier**


# AdaBoost: Step-By-Step

- **Training data**



# AdaBoost: Step-By-Step

- **Iteration 1, train weak classifier 1**



Threshold  
 $\theta^* = 0.37$

Dimension  
 $j^* = 1$


Weighted error  
 $e_t = 0.2$

Voting weight  
 $\alpha_t = 1.39$

Total error = 4

# AdaBoost: Step-By-Step

## ■ Iteration 1, recompute weights



Threshold  
 $\theta^* = 0.37$

Dimension  
 $j^* = 1$


Weighted error  
 $e_t = 0.2$

Voting weight  
 $\alpha_t = 1.39$

Total error = 4

# AdaBoost: Step-By-Step

## ■ Iteration 2, train weak classifier 2



Threshold  
 $\theta^* = 0.47$

Dimension  
 $j^* = 2$


Weighted error  
 $e_t = 0.16$

Voting weight  
 $\alpha_t = 1.69$

Total error = 5

# AdaBoost: Step-By-Step

## ■ Iteration 2, recompute weights



Threshold  
 $\theta^* = 0.47$

Dimension  
 $j^* = 2$

Weighted error  
 $e_t = 0.16$


Voting weight  
 $\alpha_t = 1.69$

Total error = 5



# AdaBoost: Step-By-Step

## ■ Iteration 3, train weak classifier 3



Threshold

$$\theta^* = 0.14$$

Dimension, sign

$$j^* = 2, \text{ neg}$$

Weighted error

$$e_t = 0.25$$


Voting weight

$$\alpha_t = 1.11$$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 3, recompute weights



Threshold

$$\theta^* = 0.14$$

Dimension, sign

$$j^* = 2, \text{ neg}$$

Weighted error

$$e_t = 0.25$$


Voting weight

$$\alpha_t = 1.11$$

Total error = 1

# AdaBoost: Step-By-Step

- **Iteration 4, train weak classifier 4**



Threshold  
 $\theta^* = 0.37$

Dimension  
 $j^* = 1$


Weighted error  
 $e_t = 0.20$

Voting weight  
 $\alpha_t = 1.40$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 4, recompute weights



Threshold  
 $\theta^* = 0.37$

Dimension  
 $j^* = 1$


Weighted error  
 $e_t = 0.20$

Voting weight  
 $\alpha_t = 1.40$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 5, train weak classifier 5



Threshold  
 $\theta^* = 0.81$

Dimension  
 $j^* = 1$


Weighted error  
 $e_t = 0.28$

Voting weight  
 $\alpha_t = 0.96$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 5, recompute weights



Threshold  
 $\theta^* = 0.81$

Dimension  
 $j^* = 1$


Weighted error  
 $e_t = 0.28$

Voting weight  
 $\alpha_t = 0.96$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 6, train weak classifier 6



Threshold  
 $\theta^* = 0.47$

Dimension  
 $j^* = 2$


Weighted error  
 $e_t = 0.29$

Voting weight  
 $\alpha_t = 0.88$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 6, recompute weights



Threshold  
 $\theta^* = 0.47$

Dimension  
 $j^* = 2$

Weighted error  
 $e_t = 0.29$


Voting weight  
 $\alpha_t = 0.88$

Total error = 1



# AdaBoost: Step-By-Step

## ■ Iteration 7, train weak classifier 7



Threshold

$$\theta^* = 0.14$$

Dimension, sign

$$j^* = 2, \text{ neg}$$

Weighted error

$$e_t = 0.29$$


Voting weight

$$\alpha_t = 0.88$$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 7, recompute weights



Threshold

$$\theta^* = 0.14$$

Dimension, sign

$$j^* = 2, \text{ neg}$$

Weighted error

$$e_t = 0.29$$


Voting weight

$$\alpha_t = 0.88$$

Total error = 1

# AdaBoost: Step-By-Step

## ■ Iteration 8, train weak classifier 8



Threshold  
 $\theta^* = 0.93$

Dimension, sign  
 $j^* = 1$ , neg


Weighted error  
 $e_t = 0.25$

Voting weight  
 $\alpha_t = 1.12$

Total error = 0

# AdaBoost: Step-By-Step

## ■ Iteration 8, recompute weights



Threshold

$$\theta^* = 0.93$$

Dimension, sign

$$j^* = 1, \text{ neg}$$

Weighted error

$$e_t = 0.25$$


Voting weight

$$\alpha_t = 1.12$$

Total error = 0

# AdaBoost: Step-By-Step

- **Final Strong Classifier**



**Total training  
error = 0**  
(Rare in practice)